

Hierarchical Context Retrieval on Knowledge Graphs Using the Tree-RAG Framework

Brij Kishor

Roll: 2411MC09

M.Tech (Mathematics and Computing)

Mini project submission Under the Supervision of

Dr. Jimson Mathew

Department of Computer Science & Engineering

Indian Institute of Technology Patna

Index:-

- ▶ Problem Statement
- ▶ Other existing RAG systems
- ▶ Graphusion: The Tree RAG Framework Architecture
- ▶ Entity and Relationships
- ▶ Generating the Knowledge Triplets
- ▶ Knowledge Graph Fusion
- ▶ Visualization: The Final Graph
- ▶ Optimizing Traversal with Filters
- ▶ Complete Workflow
- ▶ Results: Benchmark for RAG
- ▶ Key achievements
- ▶ Future enhancements

Problem Statement:-

Standard Retrieval-Augmented Generation (RAG):

It works well for simple questions on **unstructured text** (documents, articles, PDFs) but suffers for complex queries.

Fragmented Context: Standard RAG retrieves isolated text chunks , missing the full hierarchical context and parent-child relationships


Ex- "When should we use BFS?"

Low Semantic Resilience: Standard RAG fails to retrieve correct information from "noisy" or indirect queries

Ex- "Which algorithm quickly finds an element by repeatedly cutting the search region in half?"

High Latency: Slow for multi-hop questions due to sequential lookups.

Ex- "Which data structure should I use if I need the fastest lookup and I don't care about order?"



Other RAG systems:-

- ▶ **Local Graph RAG-** Knowledge is extracted into many small separate subgraphs hard to find relationships between concepts so it can't answer multi-step reasoning questions that leads to weak performance
- ▶ **Graph-RAG-** These graphs are not connected with duplicate entities, Still expensive and imperfect aggregation

Graphusion: The Tree RAG Framework

It achieves high performance by unifying three core components:

- ▶ **LLM-Driven Triplet Ingestion :** A powerful LLM automatically extracts key facts from documents.


Converts them into structured sentences called triplets, e.g., (Array , Has-property, constant-time access).

- ▶ **Embedding-Based Fuzzy Matching using HNSW :**

Queries and graph entities are converted into numerical codes called embeddings. This enables semantic matching with help of HNSW, so the system finds relevant information even without exact keyword overlap.

- ▶ **In-Memory Forest Traversal ("Instant Lookup" Phase):** Instantly retrieves full context, including ancestors ,descendants and siblings , ensuring a complete, low-latency response in such a way that No need for multiple searches.

Architecture:-



Entity and Relationships :-


- ▶ An entity is a fundamental unit of information in a knowledge graph that represents a real-world object, concept, or event, these are nodes of graph
- ▶ Relations = connections between entities that give meaning to the graph
- ▶ **Leads-to:** One concept triggers another. Example: Recursion Leads-to function calls.
- ▶ **Is-similar-to:** Entities share similarities. Example: Stack Is-similar-to Queue.
- ▶ **Has-property:** Concept has certain attributes. Example: Binary Tree Has-property two children.
- ▶ **Has-function:** Defines functional role of a concept. Example: Heap Has-function priority queues.
- ▶ **Is-implemented-by:** Concept is realized by another. Example: QuickSort Is-implemented-by recursive partitioning.
- ▶ **Depends-on:** One concept relies on another. Example: Doubly Linked List Depends-on two pointers.

Generating the Knowledge Triplets


- ▶ **Step1 -Seed Entity Extraction**
- ▶ The system uses a specialized prompt to instruct a language model to act as a "knowledge graph builder."
- ▶ This LLM is tasked with identifying key nouns or noun phrases (concepts) directly from the raw text of the document.
- ▶ This process results in a clean list of educational concepts, such as 'Array', 'Quicksort', and 'Time Complexity'.
- ▶ **Step2- Candidate Triplet Extraction :-**
- ▶ The system takes the list of extracted entities and the original text.
- ▶ It prompts the LLM again, this time to identify relationships between these entities using a strict set of predefined relations.
- ▶ The output is a collection of '(subject, predicate, object)' triplets.
- ▶ Example:(Array, Has-property, Constant-time access)

Step3-Knowledge Graph Fusion

- ▶ Goal: To transform the raw extracted triplets into a high-quality, accurate, and semantically rich knowledge graph.
- ▶ **Entity Fusion** :-Identifies and merges duplicate or semantically similar entities into a single, canonical representation.
- ▶ Example: ‘Dynamic Programming’ and ‘DP’ are fused into a single canonical entity, ensuring all related triplets point to one node.
- ▶ **Conflict Resolution** :-Resolves conflicting relationships between the same two entities. Using the original source text as a ground truth.
- ▶ Example: If the LLM generates both ‘(Array, Used-for, Sorting)’ and ‘(Array, Has-property, Sorting)’, the conflict resolution logic would analyze the text to determine the correct relationship and discard the incorrect triplet.
- ▶ **Novel Triplet Inference** :-Novel Triplet Inference enriches the knowledge graph by logically deducing new facts that weren’t explicitly written but are implied by existing relationships..
- ▶ Example: (Quicksort, Has-paradigm, Divide-and-Conquer), (Quicksort, Uses, Recursion), (Quicksort, Has-time-complexity, $O(n \log n)$)
-> (Quicksort, Is-implemented-by, recursive partitioning)




Visualization: The Final Graph



Optimizing Traversal with Filters

- ▶ Goal: When searching through a large knowledge graph, we don't want to waste time exploring parts of the graph that definitely do **not** contain the information we need.
- ▶ **Bloom Filters :-**
- ▶ **Functionality:-** A highly space-efficient filter that provides a fast, one-way check. It can definitively tell us if an entity is not present.
- ▶ **Benefit:-** If it returns 'false', we can skip that entire branch of the forest, saving significant time. If it return 'true' then it *might* be there, so we check further.
- ▶ **Cuckoo Filters :-**
- ▶ **Functionality:-** An advanced filter(works like a Bloom filter but with extra abilities) have faster lookups and support for deletions. It is designed for dynamic graphs where entities and relations are updated.
- ▶ **Benefit:-** It maintains search efficiency while allowing for a mutable dataset, which is essential for a real-world system.

Complete Workflow



- ▶ Tree RAG enables single-shot query with full hierarchical context.
- ▶ Query embedded into vector.
- ▶ FAISS IV+HNSW (Hierarchical Navigable Small World graph) finds semantic match (with filters).
- ▶ In-memory traversal retrieves ancestors + descendants .so instead of returning only one chunk, it returns the full structured context.
- ▶ Result: Complete, low-latency, context-rich response.

Results: Benchmark for Tree-RAG


Setting	T1	T2	T3	T4	T5
GPT4o zs	69.20	64.42	66.61	44.00	11.45
GPT4o RAG	64.40	65.06	69.31	40.80	10.02
GraphRAG	60.40	64.19	67.45	42.00	8.96
Ours	92.00	80.29	77.85	50.00	15.65

Evaluation on Tasks 1-5. T1, T4: accuracy; T2, T3: similarity score; T5: hit rate.


Method	NLP		CV		BIO		Overall	
	Acc	F1	Acc	F1	Acc	F1	Acc	F1
<i>Supervised Baselines</i>								
P2V[3]	0.6369	0.5961	0.7642	0.7570	0.7200	0.7367	0.7070	0.6966
BERT[7]	0.7088	0.6963	0.7572	0.7495	0.7067	0.7189	0.7242	0.7216
DeepWalk[31]	0.6292	0.5860	0.7988	0.7910	0.7911	0.8079	0.7397	0.7283
Node2vec[12]	0.6209	0.6181	0.8197	0.8172	0.7956	0.8060	0.7454	0.7471
<i>Zero-shot (zs)</i>								
LLaMA	0.6058	0.6937	0.6092	0.6989	0.6261	0.6957	0.6137	0.6961
GPT-3.5	0.6123	0.7139	0.6667	0.7271	0.6696	0.6801	0.6495	0.7070
GPT-4	0.7639	0.7946	0.7391	0.7629	0.7348	0.7737	0.7459	0.7771
<i>Zero-shot + RAG</i>								
GPT-3.5	0.7587	0.7793	0.6828	0.7123	0.6870	0.7006	0.7095	0.7307
GPT-4	0.7755	0.7958	0.7230	0.7441	0.7174	0.7200	0.7386	0.7533

Results: Benchmark for RAG

- ▶ Overall Performance: Our Tree RAG framework achieves a 2-5x reduction in latency compared to traditional methods, with sub-100 ms lookups on large-scale datasets.
- ▶ Core Innovation: By unifying embedding-based fuzzy matching with in-memory forest traversal, we enable a single, low-latency lookup that retrieves complete hierarchical context with high accuracy, even from noisy queries.
- ▶ Benefits of the Approach: This system delivers 10-20% higher recall under noisy query conditions and scales efficiently to millions of relations, making it ideal for information-rich domains.
- ▶ High-Level Summary: The Tree RAG framework successfully bridges the gap between semantic resilience and hierarchical awareness , providing a robust, high-performance solution for a new generation of knowledge-intensive RAG applications.



Gemini-2.5-Flash



Key achievements-

► Better Long Context Understanding

Advantage: Gemini 2.5 Flash handles longer educational abstracts (10,000+ characters) more effectively

Evidence: The system processes entire academic documents in chunks without losing context

Comparison: LLaMA has shorter context windows, GPT models may lose coherence in longer texts

► Superior Structured Output Generation

Advantage: Gemini consistently produces well-formatted triplets in the exact (subject, predicate, object) format required

Evidence: The prompts show strict formatting requirements that Gemini handles reliably

Comparison: LLaMA often struggles with consistent formatting, GPT-3.5/4 may hallucinate extra formatting

► Faster Processing Speed

Advantage: Gemini 2.5 Flash optimized for rapid inference

Evidence: Pipeline processes large datasets efficiently

Comparison: LLaMA requires more computational resources, GPT-4 has slower response times

► Improved Factual Accuracy in Knowledge Extraction

Advantage: Higher precision in identifying genuine relationships vs. spurious connections

Evidence: Generated triples show meaningful educational relationships

Comparison: LLaMA may hallucinate relationships, GPT can be overly creative

Future enhancements :-

- ▶ **Advanced Graph Operations:**
 - ▶ Implement graph validation and consistency checking
 - ▶ Support for graph analytics (centrality, communities, paths)
 - ▶ Enable graph versioning and diff tracking
- ▶ **Automated Quality Assurance:**
 - ▶ Add confidence scores for extracted triples
 - ▶ Implement conflict resolution mechanisms (beyond fusion step)
 - ▶ Create validation benchmarks and metrics
- ▶ **Language Extension:**
 - ▶ Complete Japanese medical domain implementation (currently incomplete)
 - ▶ Add support for other languages beyond English
 - ▶ Cross-lingual knowledge graph fusion
- ▶ **Multi-modal Integration:**
 - ▶ Extract knowledge from figures, tables, equations
 - ▶ Support PDF, HTML, and structured documents (currently only .txt)
 - ▶ Integrate with citation networks

Thank You for your
time!

